

Distributed Graph Processing:

Theoretical TLAV vs. Pregel

PETAR JOVANOVIĆ

Introduction

Distributed graph processing




Thinking Like a Vertex (TLAV) framework

- iterative execution of a user-defined vertex program over vertices of the graph
 - Vertices pass messages to adjacent vertices
- Two views: **vertices view** and **edges view**
- Synchronous (BSP)
 - Computation is based on **supersteps**, which serve as sync barriers

Pregel – most famous implementation of BSP TLAV framework


- Besides vertices and edges, it also maintains a triplets view

Theoretical TLAV

-  **Apply:** also known as vertex program, applies a user-defined **function** to each vertex in parallel; meaning that the function specifies the behavior of a single vertex v at a particular superstep S . On the first iteration, the vertex program is invoked on all vertices and the pre-defined message is passed. On subsequent iterations, the vertex program is only invoked on those vertices that receive messages.
-  **Scatter:** also known as send message, sends messages to other vertices, such that those vertices will receive the messages in the next superstep $S+1$.
-  **Gather:** receives and reads messages that are sent to a node v from the previous superstep $S - 1$ and apply the function. This function must be commutative and associative.


Pregel

Implementation of the Apply / Gather / Scatter

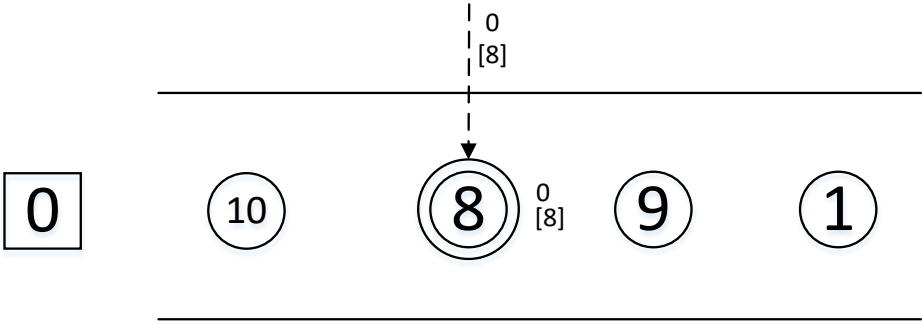
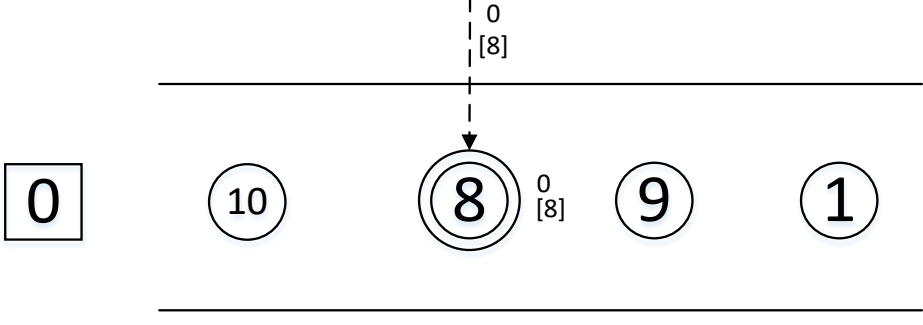
- **VProg** (corresponding to the **Apply** phase): in the case of the first superstep, the vertex value, otherwise applies the **function** over the vertex value and the received message and sends the result.
- **sendMsg** (corresponding to the **Scatter** phase): by accessing the **triplets view*** checks if it the current vertex value can **change** the destination vertex value, and if so, it sends the current value to it. Otherwise, does not send anything. 
- **merge** (corresponding to the **Gather** phase): receives messages from previous supersteps and applies the **function**.


* **Triplets view**: in addition to the vertex and edge views of the property graph, GraphX also exposes a **triplet view**. The triplet view joins the vertex and edge properties yielding an RDD that contains instances of the EdgeTriplet class. EdgeTriplet class extends the Edge class by adding the **srcAttr** and **dstAttr** members which contain the **source** and **destination** properties, respectively.

Example: Maximal value




Example: Maximal value

Theoretical TLAV algorithm	Pregel framework
	

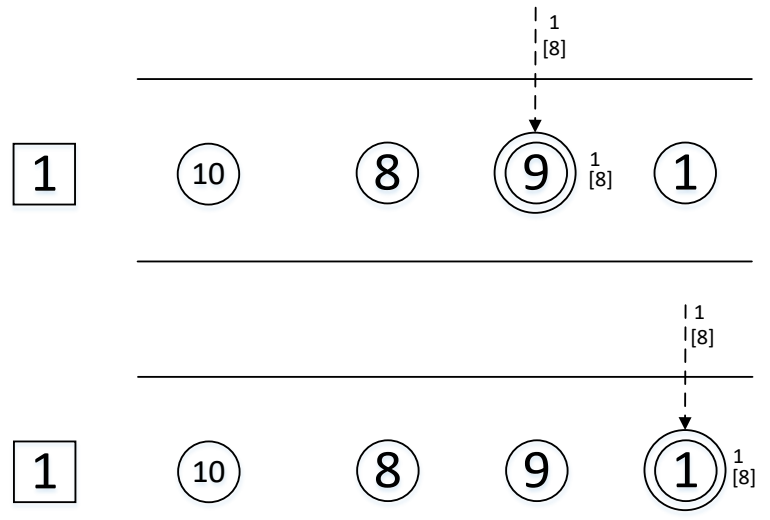
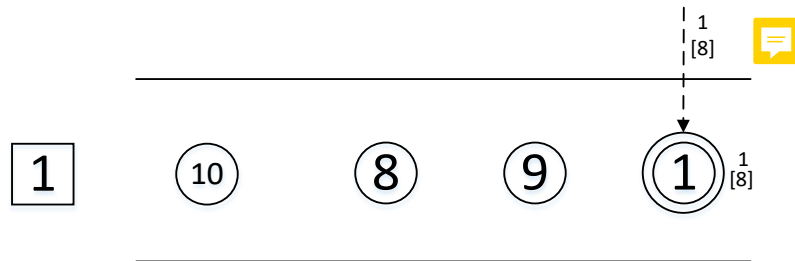



Example: Maximal value

Theoretical TLAV algorithm	Pregel framework

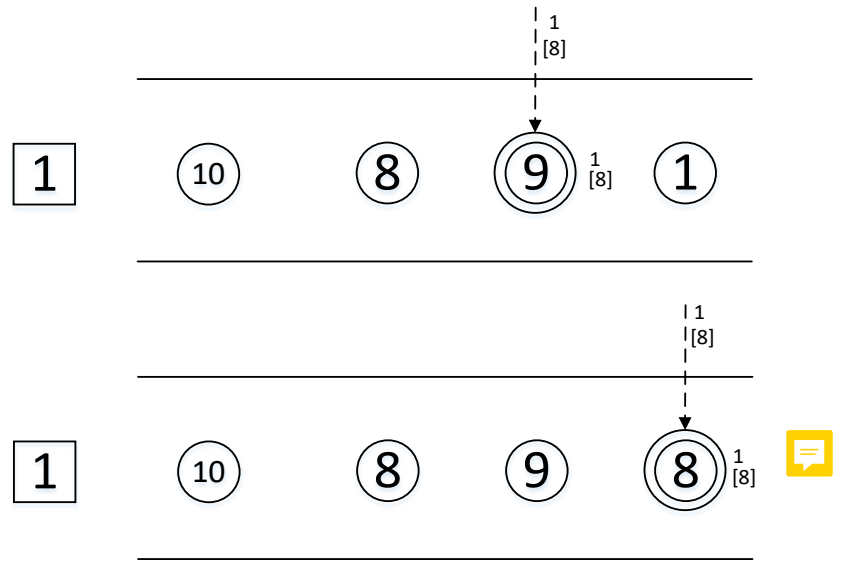
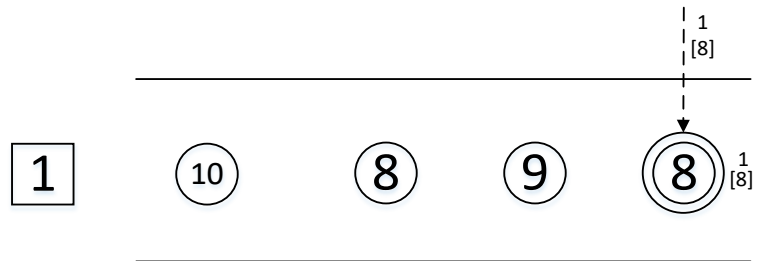



Example: Maximal value

Theoretical TLAV algorithm	Pregel framework
 <p>The diagram illustrates the Theoretical TLAV algorithm in two iterations. In the first iteration, node 9 is the maximal value with value 1 and rank 8. In the second iteration, node 1 is the maximal value with value 1 and rank 8.</p>	 <p>The diagram illustrates the Pregel framework in one iteration where node 1 is the maximal value with value 1 and rank 8.</p>

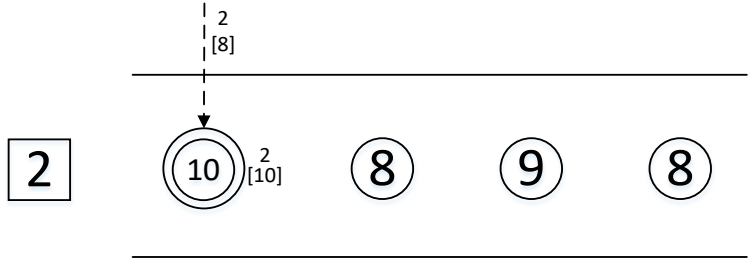
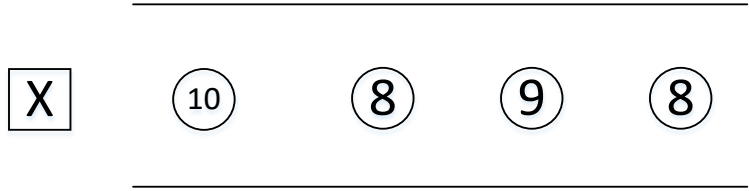



Example: Maximal value

Theoretical TLAV algorithm	Pregel framework
	



Example: Maximal value

Theoretical TLAV algorithm	Pregel framework
	




Conclusions

Pregel and GraphX implement the theoretical TLAV framework by maintaining the three views:

- **Vertices, edges, triplets** (source & destination attributes).

+ Triplets view saves us from unnecessary message sending by allowing vertices to check the destination value. 

- Maintenance of the triplets view has additional overhead, especially in the case of evolving graph topologies. 

Thank you!

Questions?